

# SQL Assignment Part 1

Devin Cline

2023-03-13

This assignment is worth 100 points. There are 15 questions (1a, 1b, 1c, 2a, etc...) and each question is worth 6.5 points each. This assignment is due Tuesday March 14 by 11:59pm CST.

Once you have completed the assignment, render the RMarkdown document to a pdf. Save the pdf as 'SQL\_Assignment\_Part1)your\_last\_name.' Upload the pdf to the SQL Assignment Part 1 link in Module 4 in Canvas.

Remember, the file will not render if the View() function is not either removed or commented out or if there are any errors in the code chunks.

Do not forget to call the sqldf library each time you invoke a new R session (that is, you exit and re-enter RStudio Cloud). The sqldf package is already installed in this project.

```
#call the sqldf package  
library(sqldf)
```

## Data Sets and Scenario

You are going to use three data tables in this assignment. These data tables are related to information about films. The name of the tables are films, reviews, and people.

Run the following code chunk below that will create the three tables in your R Studio environment.

```
#create data table films  
films <- read.csv("/cloud/project/data/films.csv")  
  
#create the data table people  
people <- read.csv("/cloud/project/data/people.csv")  
  
#create the table reviews  
reviews <- read.csv("/cloud/project/data/reviews.csv")
```

## Time to Query!

1a. Count the number of records in the people table, aliasing the result as count\_records.

Hint: The count should be 8397.

```
sqldf("SELECT COUNT(*) AS count_records  
      FROM people")
```

```
##   count_records  
## 1             8397
```

1b. Count the number of languages and countries in the films data table. Use aliases count\_languages and count\_countries, respectively.

Hint: You can use multiple COUNT() statements in a SELECT function. There should be 48 languages and 65 countries.

```
sqldf("SELECT COUNT(DISTINCT language) AS count_languages,  
        COUNT(DISTINCT country) AS count_countries  
        FROM films")
```

```
##      count_languages count_countries  
## 1                48                65
```

1c. Select the title of every film that doesn't have a budget associated with it and use the alias no\_budget\_info. Use the films data table. Limit the report to 10 rows. You can run the code chunk below to see what the first 10 rows should look like.

```
sqldf("SELECT title AS no_budget_info  
        FROM films  
        WHERE budget IS NULL LIMIT 10")
```

```
##              no_budget_info  
## 1          Pandora's Box  
## 2      The Prisoner of Zenda  
## 3          The Blue Bird  
## 4              Bambi  
## 5          State Fair  
## 6          Open Secret  
## 7      Deadline - U.S.A.  
## 8              Ordet  
## 9      The Party's Over  
## 10 The Torture Chamber of Dr. Sadism
```

2a. Generate a report that has film\_id and facebook\_likes with less than 1,000 likes from the reviews table. Limit the report to 10 rows. You can run the code chunk below to see what the first 10 rows should look like.

```
sqldf("SELECT film_id, facebook_likes  
        FROM reviews  
        WHERE facebook_likes < 1000 LIMIT 10")
```

```
##      film_id facebook_likes  
## 1         285              0  
## 2          65             491  
## 3          83             930  
## 4         111              0  
## 5          59             689  
## 6         163              0  
## 7         402              0  
## 8         251             912  
## 9         113             872  
## 10        107             975
```

2b. Count how many records have a num\_votes of at least 100,000 from the reviews table; use the alias films\_over\_100K\_votes. Hint: The count should be 1211.

```
sqldf("SELECT COUNT(*) AS films_over_100k_votes  
        FROM reviews  
        WHERE num_votes > 100000")
```

```
##      films_over_100k_votes  
## 1                1211
```

2c. Select and count the language column using the alias count\_spanish from the films data table. Apply a filter to select only Spanish from the language field.

Hint: There should 40 films in Spanish.

```
sqldf("SELECT language, COUNT(language) AS count_spanish
      FROM films
      WHERE language = 'Spanish'")
```

```
##      language count_spanish
## 1   Spanish              40
```

3. Create a report that meets the following requirements.

- Count the unique titles from the films data table and use the alias nineties\_english\_films\_for\_teens.
- Filter to include only movies with a release\_year from 1990 to 1999, inclusive.
- Add another filter narrowing your query down to English-language films.
- Add a final filter to select only films with 'G', 'PG', 'PG-13' certifications.

Hint: The count should be 310.

```
sqldf("SELECT COUNT(DISTINCT title) AS nineties_english_films_for_teens
      FROM films
      WHERE release_year BETWEEN 1990 AND 1999
      AND language = 'English'
      AND certification IN ('G', 'PG', 'PG-13')")
```

```
##      nineties_english_films_for_teens
## 1                                   310
```

4a. Calculate the average amount grossed by all films whose titles start with the letter 'A' and alias with avg\_gross\_A. Round the average to zero decimal places.

Hint: The average should be \$47,893,236 (no commas and dollar signs will be printed).

```
sqldf("SELECT ROUND(AVG(gross), 2) AS avg_gross_A
      FROM films
      WHERE title LIKE 'A%')")
```

```
##      avg_gross_A
## 1      47893236
```

4b. Select the lowest gross film in 2016 and use the alias lowest\_gross. The report should contain the title of the film.

Hint: The name of the film is Operation Chromite and the gross amount is \$31,662 (no commas or dollar signs will be printed).

```
sqldf("SELECT title, MIN(gross) AS lowest_gross
      FROM films
      WHERE release_year = 2016")
```

```
##              title lowest_gross
## 1 Operation Chromite      31662
```

5a. Select the title and duration from every film, from longest duration to shortest. Use the films data table. Limit the report to 10 rows. You can run the code chunk below to see what the first 10 rows should look like.

```
sqldf("SELECT title, durations
      FROM films
      ORDER BY durations DESC LIMIT 10")
```

##		title	durations
## 1		Carlos	334
## 2		Blood In, Blood Out	330
## 3		Heaven's Gate	325
## 4		The Legend of Suriyothai	300
## 5		Das Boot	293
## 6		Apocalypse Now	289
## 7		The Company	286
## 8		Gods and Generals	280
## 9		Gettysburg	271
## 10		Arn: The Knight Templar	270

5b. The ORDER BY can be used to sort multiple columns. It will sort by the first column specified, then sort by the next, and so on. Columns on the ORDER BY function should be separated by commas with no comma after the last column.

Create a SQL query that will generate a report exactly as shown in the code chunk below. Limit the report to 10 rows.

```
sqldf("SELECT title, release_year, durations
      FROM films
      WHERE release_year IS NOT NULL AND durations IS NOT NULL
      ORDER BY release_year, durations DESC LIMIT 10")
```

##		title	release_year	durations
## 1		Intolerance: Love's Struggle Throughout the Ages	1916	123
## 2		Over the Hill to the Poorhouse	1920	110
## 3		The Big Parade	1925	151
## 4		Metropolis	1927	145
## 5		Pandora's Box	1929	110
## 6		The Broadway Melody	1929	100
## 7		Hell's Angels	1930	96
## 8		A Farewell to Arms	1932	79
## 9		42nd Street	1933	89
## 10		She Done Him Wrong	1933	66

6a. Select the release\_year and count of films released in each year aliased as film\_count. Do not include records where release\_year is missing. Use the films data table. Limit the report to 10 rows. You can run the code chunk below to see what the first 10 rows should look like.

```
sqldf("SELECT release_year, COUNT(title) AS film_count
      FROM films
      WHERE release_year IS NOT NULL
      GROUP BY release_year LIMIT 10")
```

##	release_year	film_count
## 1	1916	1
## 2	1920	1
## 3	1925	1
## 4	1927	1
## 5	1929	2
## 6	1930	1
## 7	1932	1
## 8	1933	2
## 9	1934	1
## 10	1935	1

6b. GROUP BY becomes more useful when it is used on multiple columns or in conjunction with ORDER BY. You apply GROUP BY to multiple columns in the same manner you do ORDER BY (multiple columns separated by commas with no comma after the last column).

You want to understand budget changes throughout the years in individual countries.

Select the release\_year, country, and the maximum budget aliased as max\_budget for each year and each country. Sort the report by release\_year and country. You do not want the report to include missing values for release\_year or budget. Limit the report to 10 rows. You can run the code chunk below to see what the first 10 rows should look like.

```
sqldf("SELECT release_year, country, MAX(budget) AS max_budget
      FROM films
      WHERE release_year IS NOT NULL AND budget IS NOT NULL
      GROUP BY country, release_year
      ORDER BY release_year, country LIMIT 10")
```

```
##      release_year country max_budget
## 1          1916      USA      385907
## 2          1920      USA      100000
## 3          1925      USA      245000
## 4          1927 Germany    6000000
## 5          1929      USA      379000
## 6          1930      USA     3950000
## 7          1932      USA      800000
## 8          1933      USA      439000
## 9          1934      USA      325000
## 10         1935      USA      609000
```

7. Combining filtering and sorting provides you greater interpretability by ordering reports. You are interested in what countries have the highest average film budgets.

Create a report that meets the following requirements. a. Select the country and the average budget as average\_budget, from the films table. b. Group the results by country. c. Filter the results to countries with an average budget of more than one billion (1000000000). d. Sort by descending order of the average\_budget.

You can run the code chunk below to see what report should look like.

```
sqldf("SELECT country, AVG(budget) AS average_budget
      FROM films
      GROUP BY country
      HAVING average_budget > 1000000000
      ORDER BY average_budget DESC")
```

```
##      country average_budget
## 1 South Korea    1383960000
## 2   Hungary     1260000000
```

8. Generate a report that that returns the average budget and gross earnings for films each year after 1990 if the average gross budget is greater than 60 million.

This query is a real-world scenario. Many times, you will be asked to write a more complex query that answers a specific business question that cannot be found by playing around in applications like Excel.

To provide you some assistance, you can run the code chunk below to see what the final report should look like.

```
sqldf("SELECT country, release_year, AVG(budget) as avg_budget, AVG(gross) as avg_gross
      FROM films
      WHERE release_year > 1990
```

```
GROUP BY release_year
HAVING avg_gross > 60000000
ORDER BY avg_budget DESC")
```

##	country	release_year	avg_budget	avg_gross
## 1	USA	2016	56642742	76924036
## 2	USA	2012	41331819	62873528
## 3	USA	2015	39298329	72573303
## 4	USA	2014	35325799	62412137
## 5	USA	1992	25982030	63665195